

SY2-DATARACES: 231-250

- 231 [T/F] A variable in a thread's stack cannot be modified by the peer threads.
- 232 [T/F] A static C variable is stored in the stack of the thread that calls the function. But there will only be one instance of the variable when the function is called recursively.
- 233 [T/F] Any global variables or memory lines pointed by any global pointer are essentially memory lines that are "shared" (whose accesses need to be regulated with locks to prevent race conditions).
- 234 [T/F] When you declare "int yourArray [10]," the instruction "yourArray[1]=777" is the same as "*(yourArray+1) = 777". This is because "yourArray" is a pointer value to the base of your array and adding the pointer by "+1" will cause the value to be incremented by 4 bytes because it is an "int*" pointer.
- 235 [T/F] A cute C operation such as (e.g., i++) is always atomic (cannot be interrupted/interleaved in the middle).
- 236 [T/F] Basic machine instructions like add, sub, inc, etc. are atomic (i.e., cannot be interrupted/interleaved in the middle).
- 237 [T/F] Race condition implies that the outcome can vary from the expectation, hence is hard to debug the root cause.
- 238 [T/F] Critical section is a property that states: at most one thread runs in mutual exclusion.
- 239 [T/F] A critical section has instructions that read/write shared data, hence must be protected with locks.
- 240 [T/F] Adding too many locks or simply wrapping critical sections with locks without reorganizing your code can cause performance problems (a.k.a. the "lock step" performance).
- 241 [T/F] When using threads and locks, it's a good practice to create embarrassingly parallel tasks that do not need to synchronize too often. (See the "Parallelizing a job" slides).
- 242 [T/F] Only put variables you are protecting in the critical section. It is a good rule of thumb to make critical sections as small as you can as long as there are no synchronization issues.
- 243 [T/F] Big locks (like a single big bank lock) reduce deadlock probability and deliver better performance than small locks.
- 244 [T/F] Thread1: {lock(x); lock(y); do something ... }; Thread2: {lock(y); lock(x); do something ... }; This code will always run into a deadlock.
- 245 [T/F] If your process hits a deadlock/race condition, you can restart the process and will reproduce the bug in one run.
- 246 [T/F] If you hit a deadlock, the only way to make progress is to restart your job/app/computer, because it's impossible to undo the lock (unacquire the lock), unless you have a special OS/hardware support such as Hardware/Software Transactional Memory.
- 247 [T/F] Deadlock cannot happen if you use an odd number of locks.
- 248 [T/F] If you have CPU-intensive computations, you will get the benefit of *parallelism* (make your

program run faster) by running multiple threads on a single CPU core. (If you have a lot of laundries, you will get the benefit of parallelism by breaking them to multiple baskets but only with one washer).

249 [**T** / **F**] If you have a mix of CPU-intensive and IO-intensive (network/disk) jobs, you will get the benefit of *concurrency* by running multiple threads even on a single CPU core.

250 [**T** / **F**] — 250